

Implementation and Analysis of Various Corner Detection Algorithms

Kenichi Sakamoto^{#1}, Hyunhee Kwak^{*2}, Ziyad Kassar^{#3}

Abstract— Our study displays a comparison of various corner detection algorithms where we used a set of real world images and complexities such as noise and scale variation in order to find out which algorithm is the most efficient for feature detection and matching. The algorithms that we used were Scale Invariant Feature Transform (SIFT), Oriented FAST and Rotated BRIEF (ORB), and KAZE. Each algorithm's performance was evaluated using a set of images that featured different archeological structures such as Mt. Rushmore and Notre Dame under different conditions. Our analysis made use of the OpenCV library for the algorithm implementation in order to understand how each technique affects feature detection and matching results. We discussed the computational intensity of each of the algorithms and their applicability to real world uses in different scenarios.

Keywords— Computer Vision, SIFT, ORB, KAZE, Feature Matching, Algorithms

I. INTRODUCTION

In this project, we revisited homework assignment 4 with two main goals. Initially, in assignment 4, we used the Harris detection algorithm for finding interest points and the SIFT descriptor for feature matching. However, there are numerous different detection and descriptor algorithms that exist, such as ORB, KAZE, AKAZE, BRISK, and SIFT's own keypoint detection algorithm. Thus, we explored various different algorithms in order to understand their differences and find out which are the most efficient under different scenarios, considering that each of these detection algorithms has their own approach for finding keypoints. In this project, we opted to use those different algorithms rather than using the Harris detector in order to explore the different results for feature matching.

II. DESCRIPTION OF THE ALGORITHMS

A. SIFT

In Computer Vision, corner detection is one of the most significant keys of object detection and Harris corner detection is a popular method for identifying corner points in images. Although its

strong points such as robustness to noise and invariance to image rotation, it has its weakness on not being scale invariant.

In 2004, D.Lowe, University of British Columbia, came up with Scale Invariant Feature Transform (SIFT) in his paper, Distinctive Image Features from Scale-Invariant Keypoints, which extract keypoints and compute its descriptors. SIFT is a robust and widely used method in computer vision for detecting and describing distinctive features in images. SIFT addresses the challenge of identifying and matching image features across different scales, rotations, and viewpoints.

One of the key strengths of SIFT is to generate feature descriptors that are highly distinctive and invariant to various transformations. These descriptors encode information about the local image region surrounding each key point, capturing details such as gradient orientations and magnitudes.

SIFT employs a multi-scale approach to handle objects of different sizes within an image. By constructing a scale space pyramid and applying Gaussian blurring at different levels, SIFT can detect features at multiple scales, enabling robustness against changes in object size and viewpoint.

One drawback of SIFT is its computational complexity, especially during the keypoint detection and descriptor extraction phases. SIFT involves intensive operations such as convolutions with Gaussian kernels, gradient calculation, and histogram generation.

B. ORB

ORB(Oriented FAST and Rotated BRIEF) is a combination of FAST(features from accelerated segment test) and BRIEF(binary robust independent elementary features) descriptor algorithms with adding more features to improve the performance. It

was designed to provide a fast and efficient alternative to the SIFT.

ORB first uses the FAST algorithm for identifying keypoints in an image. Then, it applies a multi-scale image pyramid so that makes the detection algorithm scale-invariant. Then, it uses the BRIEF descriptor to find similarities. One interesting characteristic of the BRIEF descriptor is that it converts keypoints to a binary feature vector. However, this algorithm is not rotation-invariant, so ORB uses rBRIEF(rotation-aware BRIEF) to achieve rotation-invariance.

C. KAZE

KAZE algorithm is a feature detector and descriptor algorithm that uses non linear scale spaces which makes it more robust to real world image factors such as scale and noise. KAZE is different from other algorithms such as Gaussian smoothing since it does not apply linear filtering, which in return allows it to capture fine details and complex textures in the images that it is processing.

The KAZE algorithm first extracts keypoints using a determinant of the Hessian matrix method and then describes them with floating-point vectors that encode gradient magnitudes and orientations. However, because KAZE is superior to other feature descriptors in terms of robustness and fine detail capturing, it is consequently computationally intensive, so you would be sacrificing time in order to achieve much more quality feature detection.

III. CHALLENGES

Our initial challenge that we faced was that the pkl file from assignment 4 wasn't able to be unpacked, which prevented us from accessing the ground truth locations. In order to get around this issue, we had to adapt our approach by applying the ratio test method instead. Another significant challenge that we faced was that there was not much information about the KAZE algorithm on the internet, which slowed down our progress on understanding and implementing this algorithm as we did not have access to many examples of its usage. Our final challenge was the time it took to carefully understand each of the algorithms and selecting the appropriate OpenCV libraries. We had

to dive deeper into the foundations of each algorithm in order to classify it as appropriate to be used for our study.

IV. EXPERIMENTS

A. Dataset

The images we used in this project are the same as homework assignment 4. The dataset folder contains three image sets, including Episcopal Gaudi, Mount Rushmore, and Notre Dame. Each image set contains two images, A and B. They are the same picture of a building or a monument, but they have different scales or are taken from different angles.

B. Methods

We used the opencv library for implementation of the algorithms. The basic implementations we used in all three algorithms are `cv2.__.detect()` and `cv2.__.compute`. `cv2.__.detect()` function detects interest points in an image, where `__` is the name of the algorithms, as well as `cv2.__.compute` function computes the descriptors of keypoints in an image. For the feature matching algorithm, we used the Brute-Force matcher in the opencv library. By default, it applies euclidean distance when computing matching between two descriptors. However, for binary descriptors such as ORB, the hamming distance is often used. Thus, we implemented the `cv2.NORM_HAMMING` method for the ORB algorithm. For computing features across two different images, we used Lowe's ratio test. This test checks if the distance to best match is significantly smaller than the distance to the second best match by applying $\frac{\text{distance to best match}}{\text{distance to second best match}} < \text{ratio threshold}$. If this statement is true, then we store the results in a list "good_matches". Thus, we can simply score the overall ratio by calculating $\frac{\text{number of elements in good matches}}{\text{number of elements in matches}}$. We set the default threshold value to be 0.75 in our ratio test. Different implementations in each algorithm are described below.

1) SIFT

After the preparation process of the input images, such as loading, converting to grayscale using the OpenCV library, SIFT key points are detected using the ‘cv2.SIFT_create()’ function. Key points are filtered based on a threshold value to retain only the significant keypoints. To see how the accuracy is affected by this threshold, we changed the “threshold” in the range from 0 to 0.05 by 0.01 gap (0, 0.01, 0.02, 0.03, 0.04, 0.05). The threshold value beyond 0.05 is not considered because the “good match” number of one of the images reached 0m. Among the values, we picked a threshold value that shows significantly high accuracy in any of the images, and more explored keypoint numbers and accuracy at nearby numbers (0.025 and 0.018) and of the picked threshold value.

2) ORB

By default, the .create() function for ORB sets the maximum number of keypoints to be 500. We also set this number to 3000 to observe how those number of features affect the matching results. For computing the matching, brute force matching algorithm was used.

3) KAZE

Unlike the ORB algorithm, our .create() function did not set a maximum number of keypoints for KAZE, allowing it to observe a non-limited number of keypoints. The keypoints are then filtered based on a response threshold in order to eliminate weaker detections. We also used a brute force matching algorithm for the matching.

C. Results

The tables below show the different results we obtained in each algorithm with different threshold values.

1) SIFT

Thresh old value	Number of key points filtered (*ratio from previous threshold)						Accuracy					
	ND			EG			MR					
	ND_A	ND_B	EG_A	EG_B	MR_A	MR_B	Good match	Match	Ratio	Good match	Match	Ratio
0	24023	20431	73048	8948	30318	30318	4005	6637	60.34	649	4591	14.14
0.01	24023 (1)	20431 (1)	73048 (1)	8948 (1)	30318 (1)	30318 (1)	4005	6637	60.34	649	4591	14.14
0.02	15336 (.638)	14856 (.727)	59780 (.818)	4910 (.549)	19409 (.64)	19409 (.64)	3841	4622	83.10	422	2749	15.35
0.03	7064 (.461)	7845 (.528)	41131 (.688)	1807 (.368)	9425 (.486)	9425 (.486)	443	2315	19.14	269	1125	23.01
0.04	3007 (.426)	3658 (.466)	26378 (.641)	669 (.37)	4348 (.461)	4348 (.461)	180	1036	17.37	78	476	16.39
0.05	1076 (.358)	1733 (.474)	15503 (.588)	220 (.329)	1702 (.391)	4348 (.391)	32	445	7.19	17	178	9.55
										0	282	0.0

Thresh old value	Number of key points filtered (*ratio from previous threshold)						Accuracy					
	ND			EG			MR					
	ND_A	ND_B	EG_A	EG_B	MR_A	MR_B	Good match	Match	Ratio	Good match	Match	Ratio
0.025	10566	11021	49990	3011	13588	13588	1145	3321	34.48	261	1763	14.8
0.018	17628	16507	63709	5979	22256	22256	3017	5210	57.91	493	3238	15.23
										26	3377	0.77

Fig. 1. result table for SIFT

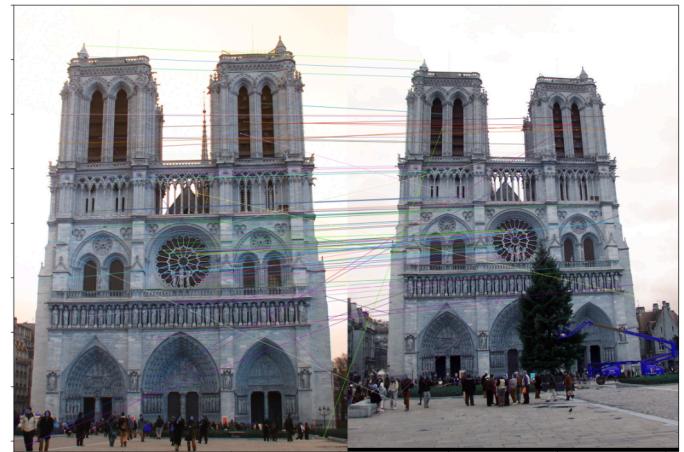


Fig. 2. result image of Notre Dame with threshold = 0.02

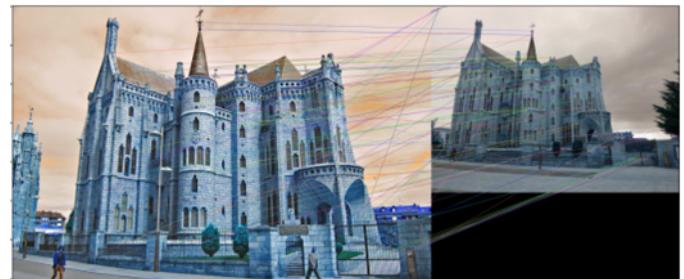


Fig. 3. result image of Episcopal Gaudi with threshold = 0.03



Fig. 4. result image of Mount Rushmore with threshold = 0.

2) ORB

Thres hold value	Number of key points filtered						Accuracy											
							ND			EG			MR					
	ND_A	ND_B	EG_A	EG_B	MR_A	MR_B	good matc h	matc h	ratio	good matc h	matc h	ratio	good matc h	matc h	Ratio	good matc h	matc h	Ratio
0	500	500	500	500	500	500	6	143	4.20	35	113	30.9 8	0	112	0			
0.001	446	482	500	171	236	236	6	137	4.38	12	62	19.3 5	0	46	0			
0.000 1	500	500	500	500	500	500	6	143	4.20	35	113	30.9 8	0	112	0			
0.000 5	500	500	500	407	500	500	6	143	4.20	30	101	29.7	0	112	0			

Fig. 5. result table for ORB with 500 features

Thres hold value	Number of key points filtered						Accuracy											
							ND			EG			MR					
	ND_A	ND_B	EG_A	EG_B	MR_A	MR_B	good matc h	matc h	ratio	good matc h	matc h	ratio	good matc h	matc h	Ratio	good matc h	matc h	Ratio
0	3000	3000	3000	3000	3000	3000	342	928	36.8 6	217	718	30.2 2	34	513	6.62 8			
0.001	1312	2095	2992	239	245	377	97	479	20.6 4	20	125	16.0 2	84	84	2.38			
0.000 1	3000	3000	3000	2726	3000	3000	342	928	36.8 6	289	675	42.8 1	34	513	6.62 8			
0.000 5	2525	2810	3000	696	2030	2030	268	817	32.8	131	279	47.0	4	375	1.07			

Fig. 6. result table for ORB with 3000 features



Fig. 7. result image of Notre Dame with threshold value = 0 in Fig. 6.



Fig. 8. result image of Episcopal Gaudi with threshold value = 0.0005 in Fig. 6.



Fig. 9. result image of Mount Rushmore with threshold value = 0 in Fig. 6.

3) KAZE

Thres hold value	Number of key points filtered						Accuracy											
							ND			EG			MR					
	ND_A	ND_B	EG_A	EG_B	MR_A	MR_B	good matc h	matc h	ratio	good matc h	matc h	ratio	good matc h	matc h	Ratio	good matc h	matc h	Ratio
0.002	4623	5372	23216	1825	14360	14360	157	3163	4.96	436	1620	26.9	0	0	0			
0.005	1274	1691	7636	467	4404	4404	219	434	50.4 6	23	233	12.4 4	0	0	0			
0.009	424	640	2818	153	1439	1439	5	157	3.18	43	81	53	33	117	28			
0.01	338	499	2236	116	1119	1119	59	117	50.4	13	60	21.6	26	88	29.5			

Fig. 10. result table for KAZE



Fig. 11. result image of Notre Dame with threshold = 0.005



Fig. 12. result image of Episcopal Gaudi with threshold value = 0.002



Fig. 13. result image of Mount Rushmore with threshold value = 0.01

D. Analysis

1) SIFT

We did not have a restricted number of features for SIFT. At the threshold of 0.03, SIFT performed best for the Episcopal Gaudi image which suggests that a moderate amount of detail was necessary in order to

capture the unique details of the building's design. At the threshold of 0.02, SIFT performed best for the Notre Dame image which indicated that this threshold was optimal for detecting significant features along with the architectural details. At the very low threshold of 0/0.01, SIFT performed best for the Mount Rushmore image which was very interesting to us and helped us conclude that this result was likely due to the less intricate and more uniform surface of the mountain in comparison to Notre Dame and the Episcopal Gaudi.

While the threshold increased by 0.01, the number of keypoints filtered decreased by a number in the range from 0.329 to 0.818 of its previous filtered keypoint sizes.

However, besides the reduction of the number of keypoints, we did not see any clear pattern in changes of the number of keypoints or in accuracy along the threshold changes. Also, the intensities of changes are diverse among images.

Although we looked into the accuracy change with nearby numbers (0.018 and 0.025) that are close to the threshold (0.02) that yielded maximum accuracy, we could not find any valid relation between threshold and the accuracy.

2) ORB

First of all, we concluded that the default number of features being 500 is too small to find the matching results. It seems that Episcopal Gaudi worked the best with this algorithm, especially with threshold value 0.0005 we got almost 50% accuracy. This proves that ORB is indeed scale-invariant as the second image of Episcopal Gaudi being much smaller than the first image. This result was interesting because the number of filtered keypoints for the second image of Episcopal Gaudi was only 696. This indicates that the majority of keypoints in there were accurate. The results for Notre Dame were acceptable but we were not able to get a single good score on Mount Rushmore. In addition, one thing we noticed

was that ORB runs noticeably faster than the other two algorithms. Despite the number of features being set 10,000 it still ran very fast.

Google colab link:
SIFT.ipynb
ORB.ipynb
KAZE.ipynb

3) KAZE

Similarly to SIFT, we did not have a restricted number of features for KAZE, but the results were still quite interesting to us. At the threshold of 0.002, KAZE performed best for the Episcopal Gaudi image which suggests that the finer details were important in this scenario. At the threshold of 0.005, KAZE performed best for Notre Dame which indicated a balance between detail and noise reduction. At the threshold of 0.01, KAZE performed best for Mount Rushmore which shows KAZE's robustness against the larger scale variations of this particular image.

V. CONCLUSION

Overall, each algorithm worked the best on different image sets. SIFT with threshold value 0.02 scored above 80% accuracy on Notre Dame. This suggests that the SIFT corner detection algorithm works better than the Harris algorithm on this image set since none of us passed 60% accuracy in homework 4 with the combination of the Harris and SIFT. For Episcopal Gaudi, we conclude that both ORB and KAZE were suitable algorithms for corner detection and matching. Those algorithms both score about 50% accuracy. KAZE also worked better on Mount Rushmore and it is the only algorithm among those algorithms that passed the accuracy above 20%. Neither SIFT nor ORB passed 10% while KAZE got about 30% with threshold value 0.01. However, since each algorithm has different strengths and weaknesses and we applied those algorithms over three different image sets, we can not simply conclude which algorithm is the best.

REFERENCES

- [1] “Introduction to SIFT (Scale-Invariant Feature Transform),” https://docs.opencv.org/4.x/d4/df5/tutorial_py_sift_intro.html (accessed May. 3, 2024)